



# BetaPoint:

## A pre-silicon performance evaluation framework

---

周耀阳

Zhou Yaoyang

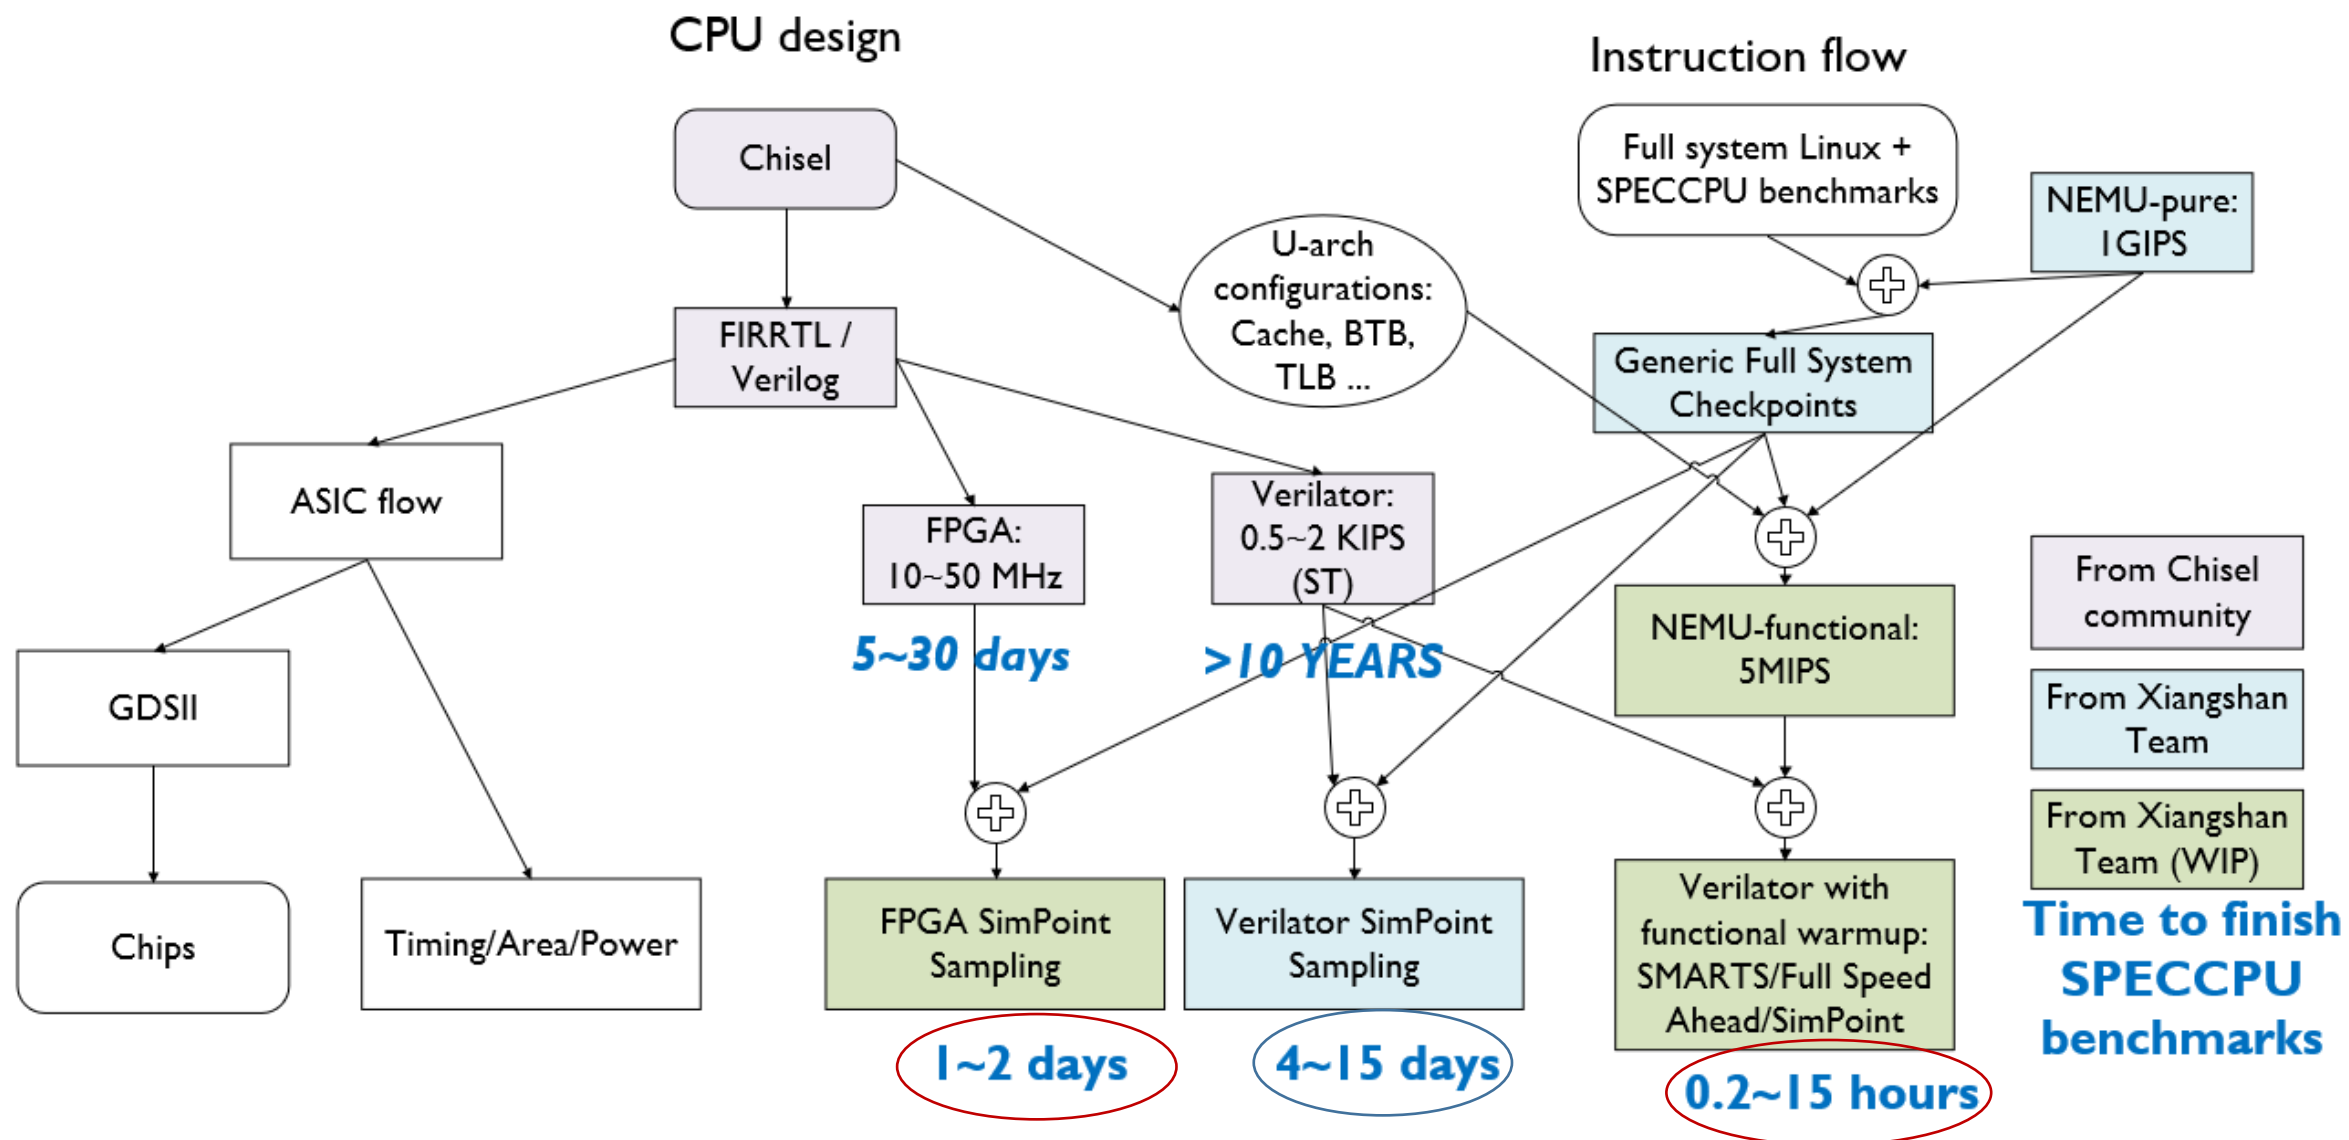
中科院计算所

ICT, CAS

2021年6月19日

June 19, 2021

# Agile performance evaluation roadmap



# ■ Agenda

- Background and motivation
  - Emerging RTL cores/SoCs
  - Lack of checkpoint and sampling support
- Current infrastructures
  - Cross-platform checkpoint format
- Ongoing works
  - Functional warmup


# ■ Opportunities and challenges

- Emerging open source cores/SoCs in RTL (Chisel, Verilog)
    - Enable agile prototyping
    - Enable researchers to produce solid results on performance, area, and timing
  - RTL emulation software is slow: **7+** years to complete SPEC2017.imagick
  - With FPGA, it is still slow: **5~30** days to finish the SPEC CPU 2006/2017 benchmarks with one FPGA.
    - **5~30 days** is still too long for performance iteration
    - Cloud FPGA cannot accommodate large cores (we use vul9p for single-core Xiangshan)
    - More FPGAs → more costs
- Can we enable sampling?

## ■ Sampling methods

Checkpoint-based sampling (SimPoint):

- Selective weighted sampling
- Large simulation points (50M~200M)



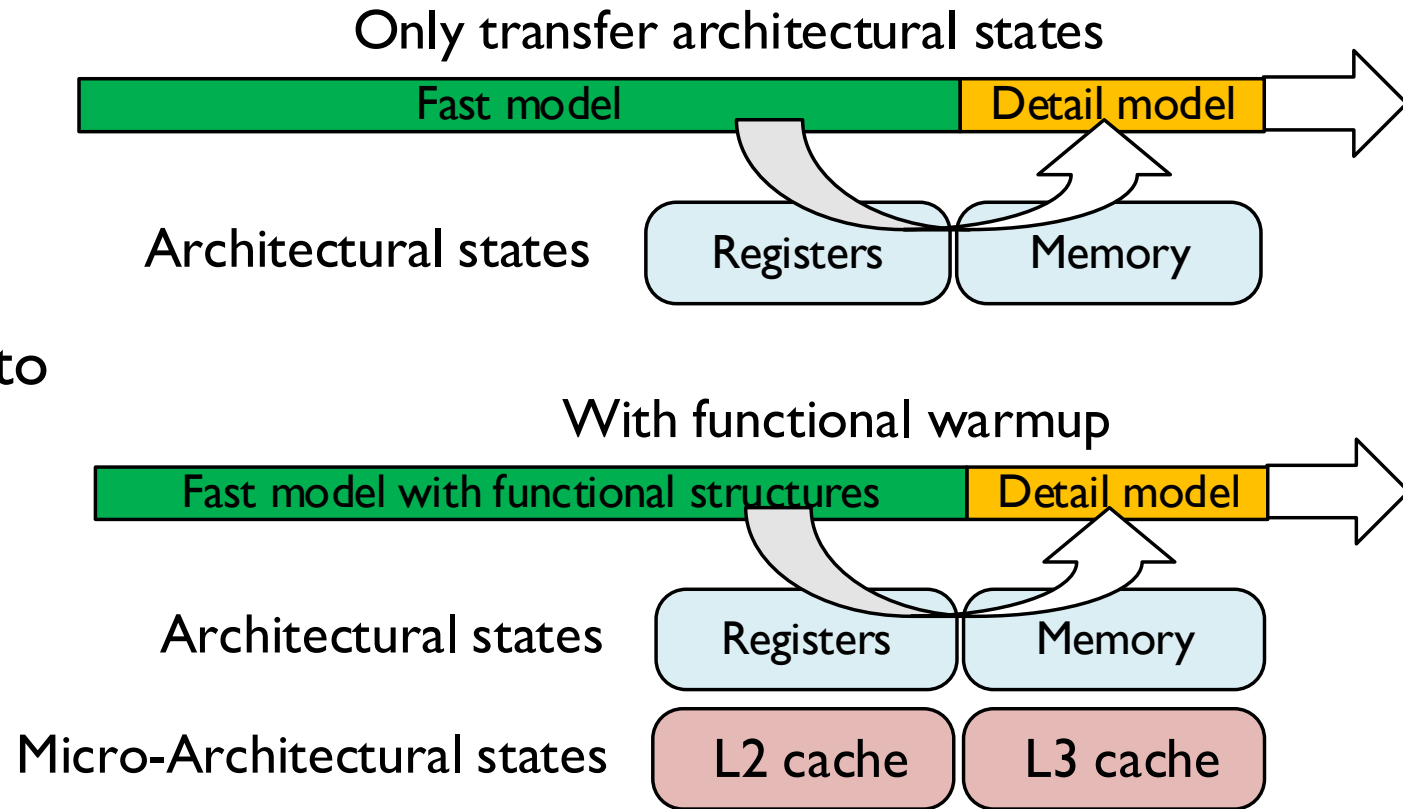
Fast-forwarding-based sampling (SMARTS):

- Uniform sampling
- Smaller simulation points (5k~50k)



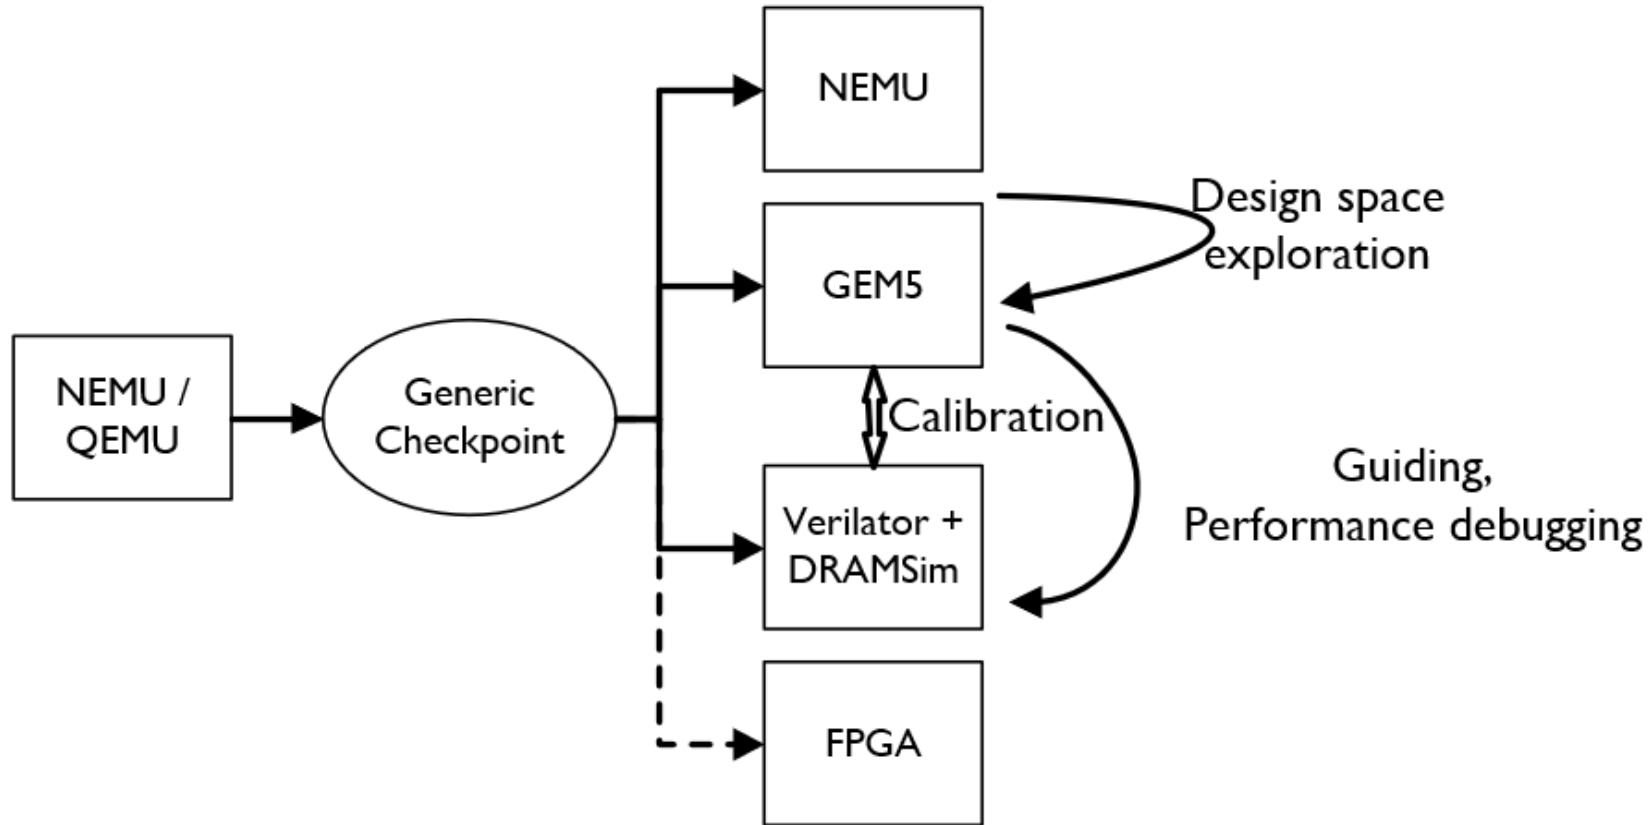
# ■ Sampling methods in research community

- Checkpoints: SimPoint
  - Requires **checkpoint support**
- Fast-forwarding: SMARTS
  - Requires **functional warmup** to speed up
- Virtualized fast-forwarding: CoolSim, DELOREAN
  - Faster than SMARTS
  - Requires (statistical) **functional warmup** to speed up



# ■ Our checkpoint format

- A generic checkpoint format (GCPT) bridges fragmented parts

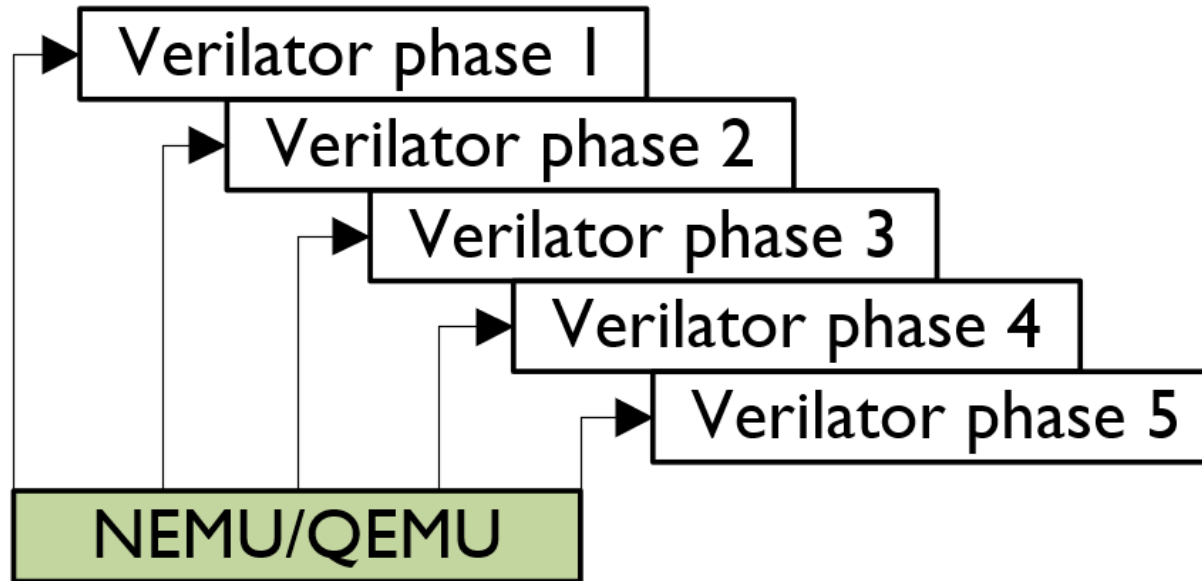


1. NEMU is a functional simulator developed by Yu Zihao, which runs at **300 MIPS~1 GIPS**.
2. GCPT is theoretically compatible with FPGA. But we have not test it on FPGA.

# ■ Current applications of NEMU+GCPT

## - Shotgun

- Shotgun method (The name is inspired by DNA sequencing)
  - Break the whole program into chunks
  - Generate checkpoints with fast models (NEMU)
  - Run detail models **in parallel**



- Reduce the time to gather “ground truth performance”
  - Fake ground truth: limited warmup length



# ■ Current applications of NEMU+GCPT

## - SimPoint on Verilator

- Performance estimation using SimPoint for Xiangshan core
- 3~14 days to finish a 100M simulation point with 1 core
  - 14 days for *mcf* because very low IPC (0.1x)

	time	ref_time	score	Coverage
astar	772.30	7020.0	9.09	0.84
mcf	686.08	9120.0	13.29	0.82
bwaves	926.09	13590.0	14.67	0.80
soplex	685.74	8340.0	12.16	0.83
povray	501.54	5320.0	10.61	0.83
dealII	644.70	11440.0	17.74	0.81
xalancbmk	645.80	6900.0	10.68	0.82
gcc	687.23	8050.0	11.71	0.81
gobmk	863.50	10490.0	12.15	0.83
h264ref	1182.14	22130.0	18.72	0.83
GemsFDTD	639.55	10610.0	16.59	0.80
zeusmp	755.00	9100.0	12.05	0.80
bzip2	1330.63	9650.0	7.25	0.82
sjeng	1169.56	12100.0	10.35	0.85
hmmer	1188.03	9330.0	7.85	0.81
namd	653.25	8020.0	12.28	0.81
gromacs	966.90	7140.0	7.38	0.81
libquantum	739.81	20720.0	28.01	0.83
perlbench	986.87	9770.0	9.90	0.82
calculix	1892.19	8250.0	4.36	0.83
tonto	980.05	9840.0	10.04	0.80
omnetpp	644.21	6250.0	9.70	0.95
sphinx3	1252.82	19490.0	15.56	0.81
milc	659.30	9180.0	13.92	0.81
lbm	676.15	13740.0	20.32	0.84
leslie3d	805.10	9400.0	11.68	0.81
cactusADM	1981.36	11950.0	6.03	0.82

# ■ Contradiction on warmup length

- Current warmup length (50M) is too short for accuracy
    - Some applications need more than 1 G warmup
    - Discussed in BLRL (The Computer Journal, 2005); Elfies (CGO, 2021)
  - Current warmup length is too long for speed
    - Two weeks to simulate 100 M instructions of mcf
- For both accuracy and speed, we must speed up warmup
- Functional warmup



中国科学院计算技术研究所  
INSTITUTE OF COMPUTING TECHNOLOGY, CHINESE ACADEMY OF SCIENCES



中国开放指令生态 (RISC-V) 联盟  
China RISC-V Alliance



鹏城实验室

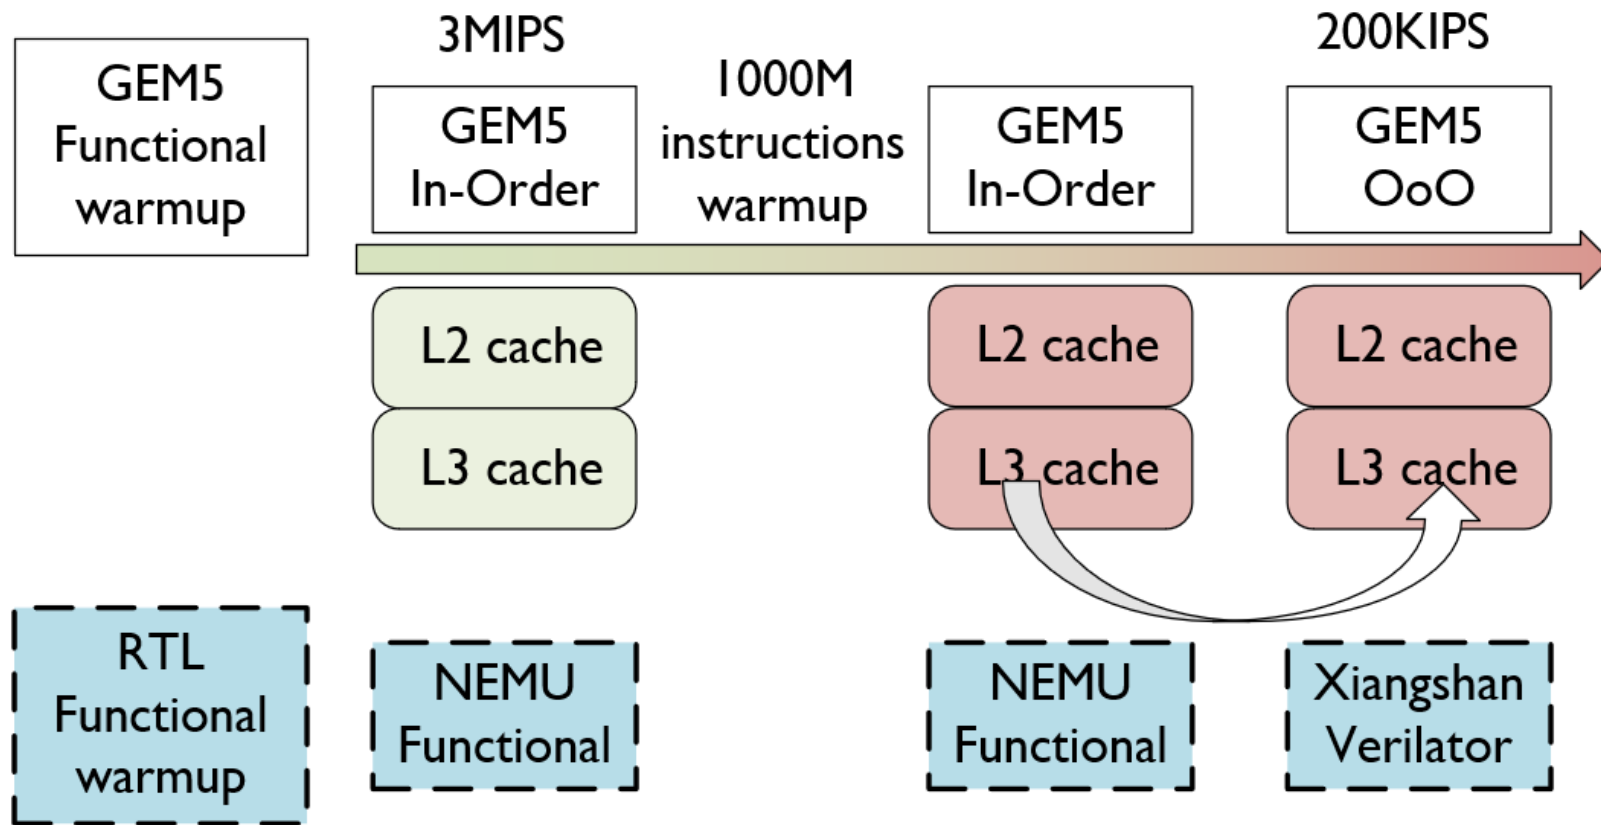
Ongoing work:

**Functional warmup for RTL emulator**


---

# ■ Functional warmup

- Warmup caches on fast model → Switch core to detailed model



# ■ Estimated speed of Verilator-FSA



	Assumed Speed	# instructions	Estimated Time
NEMU Fast-forwarding	1000 MIPS	1000M	1.0s
NEMU Functional Warmup	5.0 MIPS	1000M	200.0s
Verilator	0.5 KIPS	50K	100.0s
Total	6.64 MIPS	2000M + 50K	301.0s

If we achieve **6.64 MIPS**, we can finish SPEC CPU in **16.73h** with **100 cores**  
**1 day (12K \$)** VS. FPGA **5 days**@100MHz (**62K \$**)

# ■ Current progress

## Generating memory accesses with NEMU

- Instruction fetch addresses
- Load/Store addresses
- Functionally warm up caches @ ~16 MIPS
  - Arguably **the fastest functional warmup method** yet
  - Booting Linux in **17s**

## Comparing against existing solutions:

- GEM5 Atomic with no caches and with simple memory: < 4MIPS
- Sniper: ~2 MIPS

# ■ Next challenge: bridging C++ simulator and RTL

How to transfer **u-arch states** between C++ and RTL **Cleanly**

- Chisel BlackBox / RAM initiation?
- Verilog DPI-C?

Open problem

# ■ Towards *Asset-Light* performance evaluation

Time to evaluate/estimate SPEC CPU score:

	Without sampling		With sampling	
Devices	vu19p @ 100 MHz	Palladium Z1 @ 20 MHz	4x EPYC servers	1x <b>EPYC</b> servers
Methods	-	-	SimPoint + <b>Checkpoint</b>	Full Speed Ahead + <b>Functional warmup</b>
Pricing	~62K\$	>1550K\$	~50K\$	~ <b>12k\$</b>
Time	5 days	25 days	14 days	< <b>1 day</b> (VIP)
Scalability	At most 2-core Xiangshan	Very large SoC	Very large SoC	Very large SoC





# Thank you!

I am interested in

- Micro-architecture exploration
  - Performance modeling and evaluation
  - *Asset-Light* chip development
- (expected to graduate in 2023)

My homepage:

<https://archshine.xyz/>

